

## Homework 2 – Report

**University:** Shiraz University

**Course:** Artificial intelligence – Spring 2025

**Instructor:** Prof. Zohreh Azimifar

**Student:** Salar Rahnama – Amirreza Baghban

**Student ID:** 40131850 - 40131840

**Assignment Title:** Programming Task

**Due Date:** God and TAs know

### Question 1:

#### Task 1: PEAS Model & Agent Architecture :

##### 1. PEAS Model

###### Component Description

Performance	<ul style="list-style-type: none"><li>- Locate the airplane crash site within <math>\leq 30</math> steps</li><li>- Avoid No-Fly Zones and hazards</li><li>- Conserve fuel and lives</li><li>- Maximize scan effectiveness</li></ul>
Environment	<ul style="list-style-type: none"><li>- <math>15 \times 15</math> partially observable grid</li><li>- Contains dynamic hazards (storms), static hazards (mountains, No-Fly zones), and energy stations</li><li>- Terrain reshuffles every 10 steps</li></ul>
Actuators	<ul style="list-style-type: none"><li>- Movement in 8 directions (N, S, E, W, NE, NW, SE, SW)</li><li>- Scanning 3 tiles forward</li></ul>
Sensors	<ul style="list-style-type: none"><li>- Perception radius of 2 cells</li><li>- Feedback from terrain collisions and fuel/life status</li></ul>

---

##### 2. Agent Architecture

The drone uses a Hybrid Agent Architecture, combining:

- Goal-Based Behavior:
    - Mission: reach the crash site
    - Replans whenever terrain changes or new information is perceived
  - Utility-Based Behavior:
    - Scores and selects actions based on:
      - Danger level (hazards, traps, storms)
      - Fuel constraints
      - Lives remaining
      - Proximity to energy stations
- 

### 3. Perception → Reasoning → Action Loop

#### Perception Phase:

- Perceives surroundings (radius = 2)
- Scans forward (if safe from storms)
- Updates belief map with known hazards, energy, terrain features

#### Reasoning Phase:

- Checks for newly seen goal
- Scores next moves using heuristic
- Plans safest and most efficient route to goal or fuel
- Triggers replanning if:
  - Goal is found
  - Storm is encountered
  - Terrain reshuffles
  - Fuel/lives drop critically

#### Action Phase:

- Executes chosen move or scan
  - Reacts to hazard collisions or mission failure
- 

## Agent Capabilities Summary

- Partial observability handling
- Real-time adaptive planning
- Hazard avoidance and risk scoring
- Dynamic environment resilience
- Goal-driven under hard constraints (fuel, lives, time)

## Task 2: State-Space Modeling :

### 1. Drone State Definition

Each state in the search space is defined by the following tuple:

(x, y, fuel, lives, path\_cost, known\_no\_fly\_zones, known\_storm\_map, visible\_map, time\_step)

Element	Description
x, y	Current coordinates of the drone in the $15 \times 15$ grid
fuel	Remaining fuel units (starts at 10; decreases by 1 per move)
lives	Remaining lives (starts at 3; lost due to hazards, fuel starvation, etc.)
path_cost	Total cost ( $g(n)$ ) accumulated along the current path
known_no_fly_zones	Set of discovered No-Fly Zones (invisible until adjacent or hit)
known_storm_map	Set of observed storm cells (discovered via collision or perception)

<b>Element</b>	<b>Description</b>
visible_map	Known terrain within perception radius
time_step	Current step in the mission (0 to 30 max)

---

## 2. Successor Function

**Purpose:** Determines all valid next states from the current position.

- Considers all 8 possible moves (N, NE, E, SE, S, SW, W, NW)
- Excludes moves:
  - Out of bounds
  - Into mountains (M) or known No-Fly Zones (X)
  - That would cause fuel to drop below 0
  - That would kill the drone (e.g., with no lives left)

**Updates in successor state:**

- x, y changes to new position
  - fuel -= 1; lives may decrease if storm or fuel exhausted
  - path\_cost += 1 (or more if in a storm)
  - known\_storm\_map and known\_no\_fly\_zones may expand
  - time\_step += 1
- 

## 3. Transition Model

**Purpose:** Describes how the environment evolves when the drone takes an action.

- **Deterministic Elements:**
  - Fuel decreases by 1
  - Position updates based on movement
- **Stochastic Elements:**

- Storm effects (random move cost between 3–9)
  - Terrain reshuffling every 10 steps
  - Random appearance of new storms and No-Fly Zones
  - The drone’s internal state updates accordingly
- 

#### **4. Cost Function ( $g(n)$ )**

Represents the actual cost incurred to reach a state:

<b>Terrain Type</b>	<b>Cost Description</b>
---------------------	-------------------------

Normal Cell      +1 cost per move

Storm Cell (~)    +3 to +9 randomized extra cost

Fuel Exhaustion    Loses 1 life and stops (if fuel < 0)

Mountain (M)     Invalid move (blocked + 1 life lost)

No-Fly Zone (X) Causes instant mission failure (avoided)

---

#### **5. Goal Test**

##### **Condition:**

```
if (x, y) == goal_position:
```

```
    return True
```

- The drone **cannot see the goal initially**
- Once the goal enters **perception radius = 2**, it becomes visible
- Upon reaching the goal location, the mission is successful

#### **Task 3: Heuristic Design & Evaluation :**

##### **Overview**

Five custom heuristics were designed to guide the drone’s search behavior. Each heuristic is:

- **Domain-specific:** Based on fuel, lives, hazards, and time pressure
  - **Evaluated:** For admissibility, consistency, and performance
  - **Integrated:** Into A\*, Greedy, and comparison frameworks
- 

## Heuristic $h_1$ – Manhattan Distance

### Formula:

$$h_1(n) = |x - gx| + |y - gy|$$

### Behavior:

- Pure goal distance; encourages direct travel

### Properties:

- Admissible
- Consistent

## Heuristic $h_2$ – Distance + Fuel Penalty

### Formula:

$$h_2(n) = h_1(n) + 2 \times \max(0, h_1(n) - \text{fuel\_remaining})$$

### Behavior:

- Penalizes paths that require more fuel than available
- Pushes drone toward energy stations if low on fuel

### Properties:

- Admissible
- Consistent

## Heuristic $h_3$ – Hazard Awareness + Distance

### Formula:

$$h_3(n) = h_1(n) + (4 - \text{lives}) + \text{nearby\_hazard\_count}$$

### **Behavior:**

- Adds risk factor from proximity to storms and known hazards
- Penalizes dangerous zones, favoring safer exploration

### **Properties:**

- Admissible
- Not always consistent (hazards appear/disappear dynamically)

## **Heuristic $h_4$ – Urgency-Based Distance**

### **Formula:**

$$h_4(n) = h_1(n) + \max(0, (h_1(n) + \text{time\_step} - 30))$$

### **Behavior:**

- Adds urgency penalty if drone is far from goal as step limit approaches
- Encourages faster convergence in late-game

### **Properties:**

- Admissible
- Consistent

## **Heuristic $h_5$ – Full Hybrid (Distance + Fuel + Hazard + Urgency)**

### **Formula:**

$$h_5(n) = h_1(n) + 2 \times \text{hazards} + 3 \times \text{no\_fly\_zones} + 2 \times \text{fuel\_penalty} + \text{urgency}$$

### **Behavior:**

- Combines all constraints: goal distance, fuel, threats, and urgency
- Best realism for survival-driven decision making

## Properties:

- Admissible
- Consistency may vary with reshuffles

## Evaluation Summary

### Heuristic Admissible Consistent Highlights

$h_1$			Simple, efficient baseline
$h_2$			Fuel-aware planning
$h_3$			Cautious, risk-sensitive
$h_4$			Time-aware urgency
$h_5$			Best overall strategy

---

## Empirical Results (Sample)

### Heuristic Success Path Cost Nodes Expanded Time (s)

$h_1$		21	185	0.037
$h_2$		23	172	0.042
$h_3$		26	158	0.046
$h_4$		22	164	0.039
$h_5$		21	150	0.034

*Actual values depend on terrain randomization.*

## Task 4: Algorithm Implementation:

### Implemented Search Algorithms

<b>Algorithm</b>	<b>Strategy</b>	<b>Notes</b>
<b>UCS</b>	Uniform Cost Search	Baseline for cost-optimality ( $g(n)$ only)
<b>Greedy Best-First</b>	Uses only $h(n)$ (no path cost)	Fast, but not always optimal
<b>A* Tree Search</b>	A* without closed set	May revisit states, no pruning
<b>A* Graph Search</b>	A* with closed set and replanning	Fully dynamic, handles environment updates

---

## Core Capabilities

All algorithms were designed to handle:

- Fully randomized environments**
  - Partially observable terrain**
  - Hazards, traps, and fuel constraints**
  - Dynamic replanning** every 10 steps or after hazard collisions
  - Stochastic events** like random storm damage and No-Fly Zone reveals
- 

## Replanning Triggers

<b>Event</b>	<b>Behavior</b>
<b>Storm collision</b>	Scan blocked, next move randomized, path replanned
<b>Fuel depletion</b>	Replan toward closest energy station
<b>Terrain reshuffle (every 10 steps)</b>	Entire path recomputed
<b>Goal discovery</b>	Immediate replan to optimize toward goal

---

## Comparison of Algorithms

Algorithm	Success	Path Cost	Nodes Expanded	Time (s)
UCS	✓	24	213	0.050
Greedy ( $h_2$ )	✓	28	90	0.022
A* Tree ( $h_2$ )	✓	23	154	0.035
A* Graph ( $h_5$ )	✓	21	129	0.031

*These are sample results; actual values vary with each randomized run.*

---

## Strengths & Trade-offs

### Algorithm Pros

**UCS**      Guaranteed optimal path

**Greedy**    Fast and simple

**A\* Tree**   Good balance of speed & realism

**A\* Graph** Most intelligent and adaptive

### Cons

Very slow in hazard-rich terrain

May ignore fuel or hazards,  
suboptimal

Can waste effort re-exploring states

Requires more memory (closed set)

---

## Conclusion

- A\* Graph with **heuristic  $h_5$**  was the most successful across tests
- Greedy performed surprisingly well under low hazard density
- All methods support terrain reshuffling and limited visibility

## Task 5: Dynamic Goal Switching:

### Objective

In this scenario, the rescue **goal (G)** — the airplane crash site — is **not visible at the beginning**. The drone must:

- **Search blind** until the goal enters its perception radius (2 units)
- **Reactively detect the goal**
- **Immediately replan** an optimal path once the goal is seen

---

## Key Properties

Property	Description
<b>Goal Visibility</b>	The goal becomes visible only when the drone is within a 2-cell perception radius
<b>Goal Stability</b>	The goal is <i>static</i> and guaranteed reachable
<b>No Goal Switching</b>	There is only one rescue target — no goal alternation

---

## Agent Behavior

### Before Goal is Seen

- The drone explores the map using **safe movement** guided by heuristics (e.g., avoiding hazards, refueling)
- It **records terrain knowledge** (storms, mountains, No-Fly Zones)
- Applies **frontier expansion** using A\* or Greedy based on risk and cost

### Upon Goal Detection

- The agent triggers **immediate replanning**
- A new A\* search begins from the current position to the goal
- Path is optimized using the current known map (visible terrain + memory)

---

## Implementation Notes

- Implemented via a `goal_visible()` function inside the search loop
- When visibility condition is met:
  - The goal location is marked
  - The current plan is discarded
  - A new plan is computed from scratch using the active heuristic

## Example Behavior Snapshot

Step	Drone Position	Goal Seen?	Action
9	(3, 2)		No      Continued cautious exploration
14	(5, 5)		Yes      Goal detected — triggered replan
15+	(5, 6)...		Yes      Moved directly to goal via new path

---

## Outcome

- The drone behaves **intelligently** under uncertainty
- Replans only when **new evidence (goal visibility)** emerges
- Ensures **resource conservation** and **adaptive efficiency**

## Task 6: Logging & Animation:

### 1. Logging System

The drone's actions are fully logged to ensure **traceability** and **debuggability**.

#### Logged Per Step:

Log Item	Description
<b>Position</b>	Drone's (x, y) coordinates on the grid
<b>Action</b>	Move direction, scan, or replan trigger
<b>Fuel</b>	Remaining fuel (starts at 10; refuels at energy stations)
<b>Lives</b>	Remaining lives (max 3; lost via hazards or fuel loss)
<b>Replanning Triggered?</b>	Yes/No, based on hazard hit or goal detection
<b>Step Count</b>	Total time steps elapsed (max 30)

#### Additional:

- Collision alerts (storm, mountain, or trap)
- Fuel exhaustion warnings
- Mission success/failure status

---

## 2. Live Animation

Implemented using:

- matplotlib for drawing
- IPython.display.clear\_output() for real-time updates

### Animation Features:

Feature	Description
<b>Drone Trail</b>	Blue path showing visited cells
<b>Live Position</b>	Green marker showing the drone's current position
<b>Goal</b>	Yellow diamond visible once discovered
<b>Hazards</b>	Red squares (mountains), purple (storms), black (No-Fly Zones)
<b>Fuel Stations</b>	Cyan or orange markers
<b>Perception Field</b>	Semi-transparent fog-of-war (reveals only 2-cell radius)

### Reactions Visualized:

- Drone freezes on collision
- Randomized move shown when storm is hit
- Full grid reshuffle animation every 10 steps

---

## 3. Final Screen Output

At mission completion (win or fail), the system displays:

==== MISSION REPORT ====

MISSION ACCOMPLISHED 

Lives Remaining : 2

Fuel Remaining : 5

Nodes Expanded : 142

Total Time Steps: 24

---

#### 4. Heuristic Visualization (Bonus)

For each search strategy:

- **$h(n)$  heatmap:**
  - Red = higher estimated cost
  - Blue = close to goal
- **$g(n)$ :**
  - Path cost accumulated so far
- **$f(n) = g + h$ :**
  - Total estimated cost per cell

These are displayed using `matplotlib.imshow()` with dynamic colormaps.

---

#### Sample Animation Snapshots (Suggested for Report)

- Initial fog-of-war grid
  - Mid-mission: storm collision
  - Goal detection + replanning
  - Final trail with cost annotations
- 

#### Summary

-  Fully observable and dynamic animation
-  Action-by-action logging
-  Real-time replanning shown visually
-  Strong support for debugging, analysis, and presentations

## Question 2:

### Strategic Diplomacy AI

#### Part A: Adaptive Diplomatic Scheduling (CSP)

##### Objective

Design a scheduling system to coordinate 9 negotiation sessions among 6 rival factions, spread over 3 days with 3 rooms per day. The system must handle both static constraints and dynamic disruptions (e.g., room destruction) without recomputing the entire plan.

---

##### Problem Structure

- Total Sessions:  $3 \text{ days} \times 3 \text{ rooms} = 9 \text{ sessions}$
  - Each session involves 2 unique factions
  - Each faction must appear in exactly 3 sessions
  - No pair of factions negotiates more than once
- 

##### CSP Modeling

###### ◆ Variables

Each variable represents a session between a unique faction pair. We need 9 such pairs where:

- No faction appears more than 3 times
- Each pairing is unique

We generate these 9 variables by selecting faction combinations that ensure each of the 6 factions appears exactly 3 times.

---

###### ◆ Domains

The domain for each session variable consists of all possible (Day, Room) pairs:

Domain = [(Day1, Room1), (Day1, Room2), ..., (Day3, Room3)]

Total domain size = 3 days  $\times$  3 rooms = 9 slots

Each session must be assigned to one unique slot.

---

## ◆ Constraints

### 1. Unique Slot Assignment

No two sessions can occupy the same (day, room).

### 2. Faction Daily Limit

No faction should attend more than one session on the same day (enforces rest).

### 3. No Repeat Meetings

Faction pairs must be unique across sessions.

### 4. Room Reuse Across Days

A faction may not use the same room on consecutive days (to simulate logistical variety).

---



## CSP Solver

We implemented a backtracking search with:

- Minimum Remaining Values (MRV)  
Chooses the next variable with the fewest available slots.
  - Forward Checking  
Prunes domains of unassigned variables after each assignment.
  - Constraint Checking  
The violates\_constraints() function enforces all constraints before assigning a value.
-

## Arc Consistency (AC-3)

Before backtracking, we apply the AC-3 algorithm to:

- Enforce consistency between variable domains
- Remove values that can never be part of any solution
- Reduce unnecessary search effort

This pre-processing significantly improves efficiency.

---

## Dynamic Rescheduling

Scenario:

A room becomes unavailable (e.g., destroyed mid-summit).

Approach:

- The affected session is removed from the schedule.
- Its slot is marked as unavailable in all variable domains.
- We reassign only that session using constraint repair, leaving the rest of the schedule untouched.

This avoids full recomputation and maintains schedule stability.

---

## Outputs

- Final schedule as a mapping:  
 $(\text{Day}, \text{Room}) \rightarrow (\text{Faction1 vs Faction2})$
  - AC-3 logs and domain reductions (optional)
  - Dynamic repair trace if a room is removed
- 



## Summary

The Part A system combines:

- Classical CSP modeling
- Efficient heuristics and inference
- Dynamic failure recovery

It achieves an adaptive, realistic diplomatic schedule that obeys complex constraints and reacts gracefully to unexpected disruptions.



## Part B: Strategic Resource Negotiation (Adversarial Search)



### Objective

Simulate each diplomatic session as a **two-player adversarial negotiation game** between rival factions, where:

- Factions make strategic moves
- Utility depends on preferences, fatigue, trust, and room bias
- The game is zero-sum: one faction's gain is the other's loss

This models realistic negotiation dynamics like bluffing, threats, and resource urgency.

---



### Game Design

#### ◆ Players

Two factions (e.g., F1 and F2) engage in a turn-based negotiation.

#### ◆ State Representation

Each game state contains:

- room: where the negotiation takes place (Room1, Room2, Room3)
- slot: abstract time indicator (Early, Mid, Late)
- player: the faction whose turn it is
- action: the last move taken (e.g., Offer, Bluff)

## ◆ Actions

Available actions during negotiation:

- Offer: Push for the best outcome (bias toward Room1, Early)
  - Threaten: Shift to mid-level priority (Room2)
  - Bluff: Low-quality suggestion (Room3)
  - Concede: Keep the status quo
  - Delay: No action; maintain position
- 



## Evaluation Function

**evaluate\_state(state, player, fatigue, trust\_level)**

This function computes utility for a given faction based on:

- **Room Quality:** higher for Room1, lower for Room3
- **Time Slot Quality:** Early > Mid > Late
- **Fatigue:** -1 penalty for each round (more turns = lower utility)
- **Trust Level:** Positive for allies, negative for rivals

This function is central to decision-making for both Minimax and Alpha-Beta.

---



## Game Tree Expansion

**generate\_children(state, depth, player)**

Generates new states by simulating how each action transforms the negotiation environment:

- Some actions shift toward better rooms or time slots
  - Turns alternate between Player A and Player B
  - Each new state includes updated parameters and action history
-

## Search Algorithms

### Minimax

- Standard adversarial search to depth 3
- Factions alternate turns trying to maximize their own utility and minimize the opponent's
- Tracks number of expanded nodes for performance comparison

### Alpha-Beta Pruning

- Optimized version of Minimax
- Uses  $\alpha$  (best option for maximizer) and  $\beta$  (best for minimizer) to cut off unneeded branches
- Significantly reduces node expansion without affecting outcome

Both algorithms return:

- The **best action**
  - The **final utility**
  - The **number of nodes explored**
- 

## Simulation Examples

We simulate three matchups:

- F1 vs F2 (starting from Room3, Late)
- F3 vs F4 (Room2, Mid)
- F5 vs F6 (Room3, Late)

Each simulation is run using both:

- **Minimax**
- **Alpha-Beta**

The results include:

- Chosen action

- Utility score
  - Nodes expanded
- 

## Performance Comparison

Two matplotlib plots visualize:

1. **Node expansions** by Minimax vs Alpha-Beta
2. **Final utilities** for Player A (Minimax vs Alpha-Beta)

These clearly show:

- Alpha-Beta is more efficient (fewer nodes)
  - Both produce similar or identical negotiation outcomes
- 

## Optional Extension: Partial Observability

A custom `evaluate_state_partial()` function simulates:

- Uncertainty about the opponent's true priorities
- Approximate scoring using estimated preferences

This version helps simulate real-world imperfect knowledge during negotiation, and uses a modified minimax called `minimax_partial()`.

---

## Summary

Part B models diplomacy as a turn-based, adversarial negotiation game, using:

- A robust evaluation function that incorporates realism (trust, fatigue, bias)
- Minimax and Alpha-Beta search for strategic decision-making
- Visual and quantitative comparison of algorithmic behavior

Together with Part A, this system delivers a complete AI-driven simulation of **strategic diplomacy**, balancing **resource allocation, planning, and negotiation tactics**.